Language Modeling

CMSC 473/673 Spring 2017

Bryan Wilkinson

Discussion of Kilgarriff Paper

- Does anyone disagree with Kilgarriff?
- What are the alternatives?
  - UkWaC and WaCkypedia
  - UMBC WebBase

Language Modeling

- A Language Model is a collection of probabilities of some unit of language
  - Word and character level models are the most common.
- Why do we need to know the probability of language?
  - Which spelling is more likely?
  - Which verb is more likely?
- What is the probability of language?
  - Of a word?
  - Of a sentence?

Counting Words

- Before we can calculate probability, we need to count things.
  - Seems like a simple task but,
- What is a word?
  - Still an open question, but mostly agreed upon
  - Is *a lot* one word or two?
  - How about *post office*?
- English is easy to count
  - What about a language with more complex morphology (We will talk more about this later)
  - Iyewičhamačheča (Lakota “I resemble you” or “you resemble me”)
- Other minor decisions
  - Does capitalization matter?
  - How about plurals? Or other endings?

Probability of a Word

- Counting is the hard part

$$P(w) = \frac{C(w)}{N}$$

In Class Activity

Markov Assumption

- The future doesn't depend on the past
  - But the future **does** depend on the past, so we will just limit the amount of past we look at
- In NLP this means when calculating the probability of a word given a context, we often just look at the preceding one or two words. (or none!)
- In terms of probability this looks like
  - P(*homework* | *the dog ate my*) ≈ P(*homework*)
  - P(*homework* | *the dog ate my*) ≈ P(*homework* | *my*)
  - P(*homework* | *the dog ate my*) ≈ P(*homework* | *ate my*)
- We call these small chunks of words N-grams
  - Unigram = homework
  - Bigram = my homework
  - Trigram = ate my homework
- What are all the bigrams and trigrams in the sentence "the dog ate my homework"

Probability of an N-Gram

- Calculating probability of an N-Gram is almost as easy as a word
- Technically we are estimating using Maximum Likelihood Estimation

$$P(w|w_{n-1}) = \frac{C(w_{n-1}w_n)}{\sum_w C(w_{n-1}w)}$$

$$P(w|w_{n-1}) = \frac{C(w_{n-1}w_n)}{C(w_{n-1})}$$

Probability of a Sentence

- To calculate the probability of a sentence, we use the chain rule

$$P(w_1 w_2 w_3 ... w_n) = P(w_1)P(w_2|w_1)P(w_3|w_1, w_2)...P(w_n|w_1, w_2, w_3 ... w_{n-1})$$

- P(Today is the second day of class) = P(Today | <start>) * P(is | Today <start>) * P(the | is Today <start>) * P(second | the is Today <start>) * P(day | second the is Today <start>) * P(of | day second the is Today <start>) * P(class | of day second the is Today <start>) * P(<end> | class of day second the is Today <start>)

Probability of a Sentence

- To calculate the probability of a sentence, we use the chain rule

$$P(w_1 w_2 w_3 ... w_n) = P(w_1)P(w_2|w_1)P(w_3|w_2)...P(w_n|w_{n-1})$$

- P(Today is the second day of class) ≈ P(Today | <start>) * P(is | Today) * P(the | is) * P(second | the) * P(day | second) * P(of | day) * P(class | of) * P(<end> | class)

Probability of a Sentence

- The following counts are from Google Books N-Gram Corpus v2
  - Via corpus.byu.edu

C(<start>,Today)	
C(Today, is)	49215
C(is, the)	101249451
C(the, second)	17216997
C(second, day)	396983
C(day, of)	5639617
C(of, class)	672495
C(class,<end>)	2703083

C(<start>)	
C(Today)	3251311
C(is)	1368855691
C(the)	7726878625
C(second)	42738935
C(day)	78211120
C(of)	5035745089
C(class)	27260152
C(<end>)	7805955213

Probability of a Sentence

P(Today | <start>)	
P(is | Today)	
P(the | is)	
P(second | the)	
P(day | second)	
P(of | day)	
P(class | of)	
P(<end> | class)	

Probability of a Sentence

- What additional data would we need to calculate the probability of :
  - Today is the eighth day of class ?
  - Today is the 57th day of class ?
  - Today is the somethingth day of class?

Evaluating Language Models

- LMs give us probabilities of a sentence
  - How do we know if they are working well
  - What is the “correct” probability of a given sentence
- We could evaluate them using some task that has LMs as part of it
  - Test a machine translation, speech recognition, etc system and see which LM does better
  - Can conflates any errors
- So we try to evaluate LMs as a standalone system
  - This is hard
  - We do this using perplexity, a measure from information theory.

Perplexity

- Wikipedia (and probably statistics books) define perplexity as 2 raised to the entropy of the distribution
  - Technically correct, but not very intuitive (to me at least)
- Jurafsky gives the definition as the probability an LM gives to a test set, normalized by the number of words in that test set.
  - The idea is that words and sequences of words that actually exist in text should have a relatively high probability
- The assumption is that perplexity should correlate to performance on the end task
  - Often is the case!
  - But not always!

$$\sqrt[N]{\prod_{i=1}^{N} \frac{1}{P(w_i|w_1 \ldots w_{i-1})}}$$

Perplexity Examples

- This data comes from a system I am working on that has a perplexity of about 3

Input	Output
Who are you ?	Who are you ?
This caused problem like the formation and growth of slums which most of the time is not safe due to their unhealthy environment .	This caused problem like the both and growth of of which most of the time not not not be to their . .
Tokyo is often referred to as a city, but is officially known and governed as a "metropolitan prefecture" .	New is often referred to as a city , but is officially known and workforce as a " the prefecture " .
The Tokyo metropolitan government administers the 23 Special Wards of Tokyo	The financially of government classified the 00 Special Wards of New

Character Level Language Models

- What are they useful for?
  - Language ID
  - Spell Check
- They are smaller
  - Just need to count letters and a few other symbols
- Becoming more popular with Neural Networks
  - The Unreasonable Effectiveness of Recurrent Neural Networks
- Do they represent actual linguistic units?

Character LM Example

- Let’s say we want to build a simple language ID system
- Bigram model over 2GB of wackypeadia produces the following counts

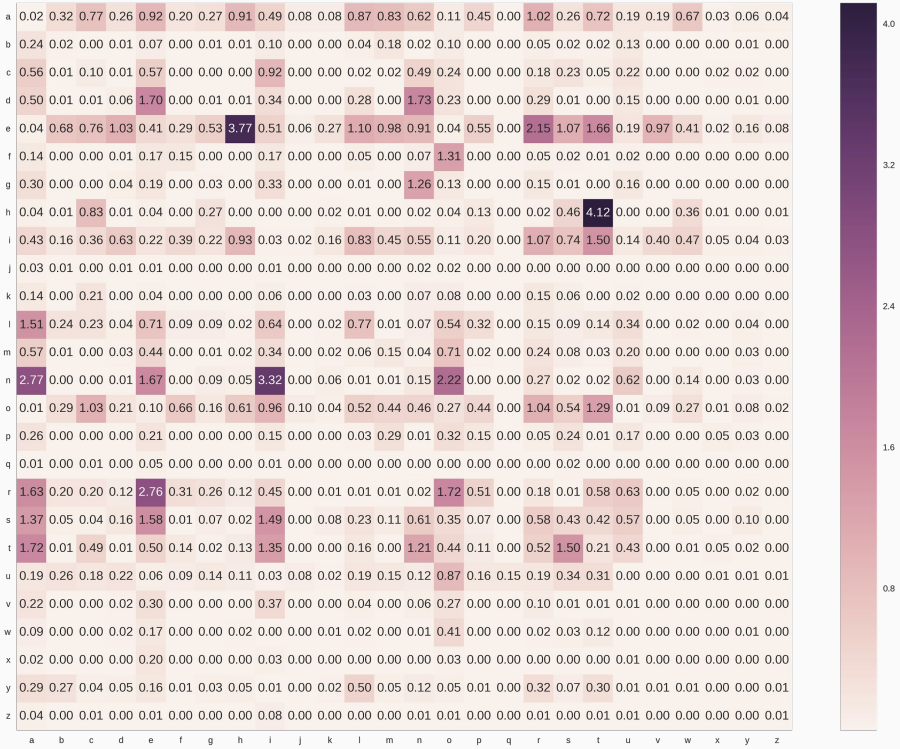

Character LM Example

- Let’s say we want to build a simple language ID system
- Bigram model over 2GB of French UN text produces the following counts

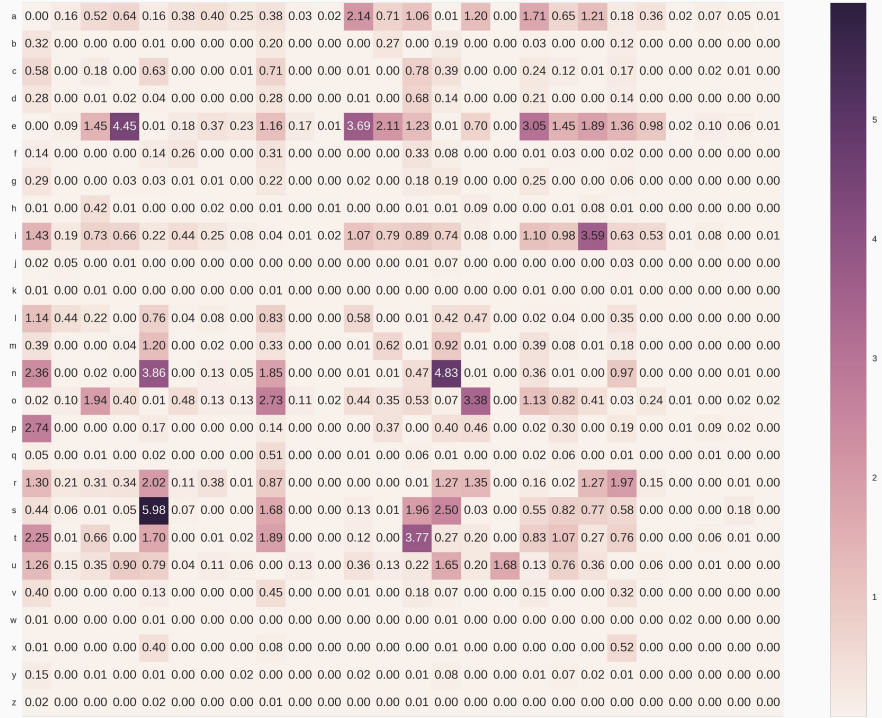

Character LM Example

- Let’s say we want to build a simple language ID system
- Bigram model over 2GB of Spanish UN text produces the following counts

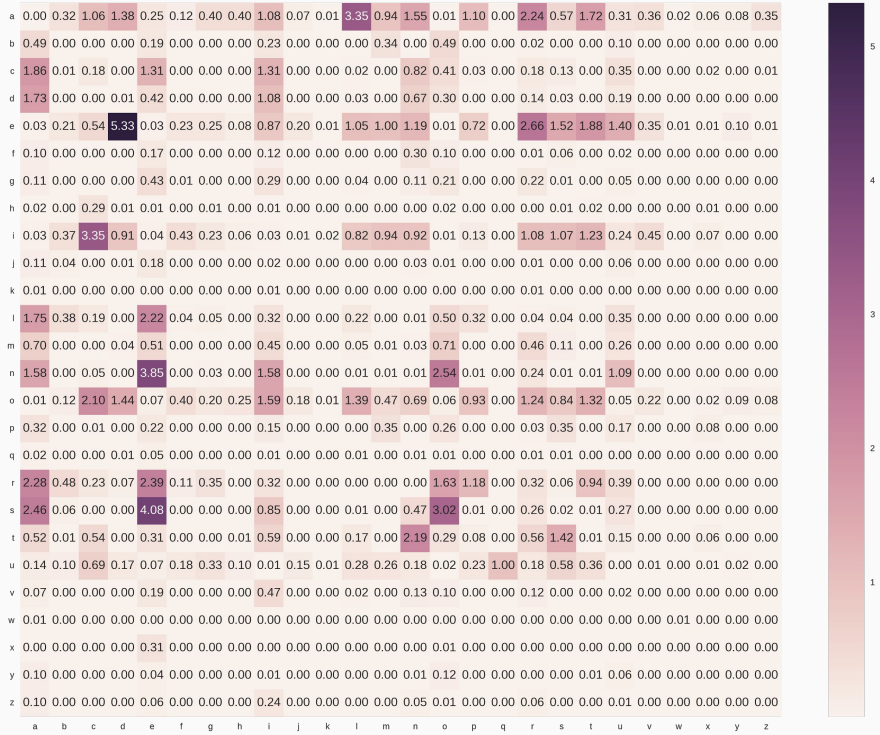

Character LM Experiment

- Let’s try the following sentences and see the results
  - Today is the second day of class
  - La universidad tienen 50 años.
  - Marius se sentit fier de cet inconnu.
  - Melissa Villaseñor is on SNL
- What happens when we put in a language we don’t have a model for
  - What can this tell us?

Reminder

- HW 0 “Due” Tuesday